

Project #4: Multi-Core Performance Evaluation & Analysis

18-640, Foundations of Computer Architecture,
Fall 2015

Outline

- What is a compute stick?
- How to use a compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

Outline

- What is a compute stick?
- How to use a compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

Intel's Compute Stick



Peripherals:

- USB 2.0 port
- Power port
- HDMI
- Micro SD Card (not needed for our purpose)

Configuration:

- Quad-core Intel Atom Processor
- Windows 8.1
- 32 GB storage
- 2 GB memory
- Wifi card

How to use it?

- HDMI to DVI Adaptor (optional)



- USB Hub for external keyboard and mouse




Outline

- What is a compute stick?
- How to use compute stick?
- **Multi-Core Execution Libraries**
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

OpenMP

- Stands for Open specifications for MultiProcessing
- First released in 1997 for Fortran; released for C++ the following year
- Supports thread-based parallelism
- Uses fork-join model (may or may not support nested parallelism):



- May differ for each vendor
- Code simplicity, modularity, and scalability over Do-It-Yourself multithreads

Example: Compute the nth fibonacci number

Pthread version:

```
int fib(int n)
{
    if (n < 2) return n;
    else {
        int x = fib(n-1);
        int y = fib(n-2);
        return x + y;
    }
}

typedef struct {
    int input;
    int output;
} thread_args;

void *thread_func ( void *ptr )
{
    int i = ((thread_args *) ptr)->input;
    ((thread_args *) ptr)->output = fib(i);
    return NULL;
}

int main(int argc, char *argv[])
{
    pthread_t thread;
    thread_args args;
    int status, result;
    int thread_result;
    int n = atoi(argv[1]);
    if (n < 30) result = fib(n);
    else {
        args.input = n-1;
        status = pthread_create(&thread, NULL, thread_func, (void*) &args);
        result = fib(n-2);
        pthread_join(thread, NULL);
        result += args.output;
    }
    printf("Fibonacci of %d is %d.\n", n, result);
    return 0;
}
```

OpenMP version:

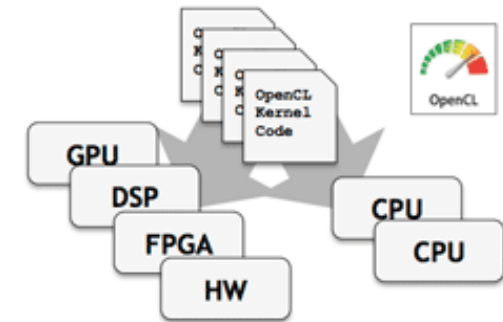
```
int fib(int n)
{
    if (n < 2) return n;
    else {
        int x = fib(n-1);
        int y = fib(n-2);
        return x + y;
    }
}

int main(int argc, char *argv[])
{
    int result;
    int n = atoi(argv[1]);
    if (n < 30) result = fib(n);
    else {
        int i,j;
        #pragma omp task shared(i)
        i = fib(n-1);
        #pragma omp task shared(j)
        j = fib(n-2);
        #pragma omp taskwait
        result = i + j;
    }
    printf("Fibonacci of %d is %d.\n", n, result);
    return 0;
}
```

Advantage of OpenMP: Easy to read, modular, sclable; also less overhead.

OpenCL

- Publically released in December, 2008 by Khronos Group
- Open standard for developing cross-platform, vendor agnostic, parallel programs that run on current and future multi-core processors



- In comparison to OpenMP, OpenCL will have some extra overhead of compiling the kernel at runtime

Outline

- What is a compute stick?
- How to use compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- Questions?

OpenCV

- Open source BSD licensed computer vision library
- Started by Intel Research in 1998
- Written in C++ with bindings available for Python, Java, Matlab
- Some questions in Project 4 will be based on image processing

Image: How is it represented?


- Two dimensional array of pixels
- Binary image:
 - Pixels are bits
- Grayscale image:
 - Pixels are scalars
- Color image:
 - Pixels are vectors

Index	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	10
1	11	12	13	14	15	16	17	18	19	20
2	21	22	23	24	25	26	27	28	29	30
3	31	32	33	34	35	36	37	38	39	40
4	41	42	43	44	45	46	47	48	49	50
5	51	52	53	54	55	56	57	58	59	60
6	61	62	63	64	65	66	67	68	69	70
7	71	72	73	74	75	76	77	78	79	80
8	81	82	83	84	85	86	87	88	89	90
9	91	92	93	94	95	96	97	98	99	100

Grayscale Image in form of a matrix


Using OpenCV for Edge Detection

- Input image:
- Canny algorithm: for edge detection



Using OpenCV for Line Detection

- Input image:
- Canny + Hough transform: edge-detection and feature extraction



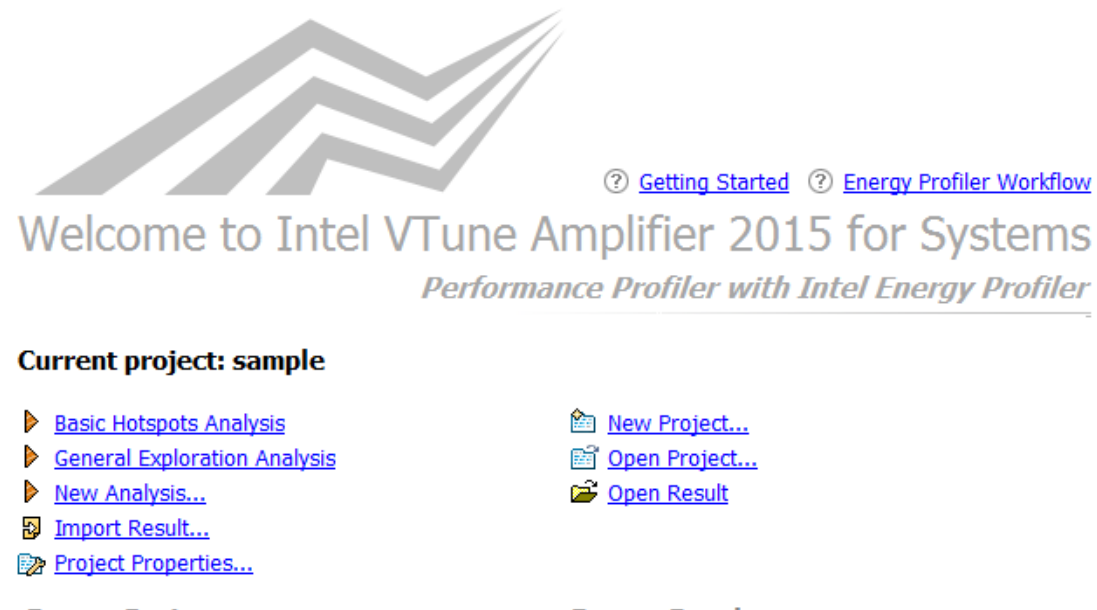
Outline

- What is a compute stick?
- How to use compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

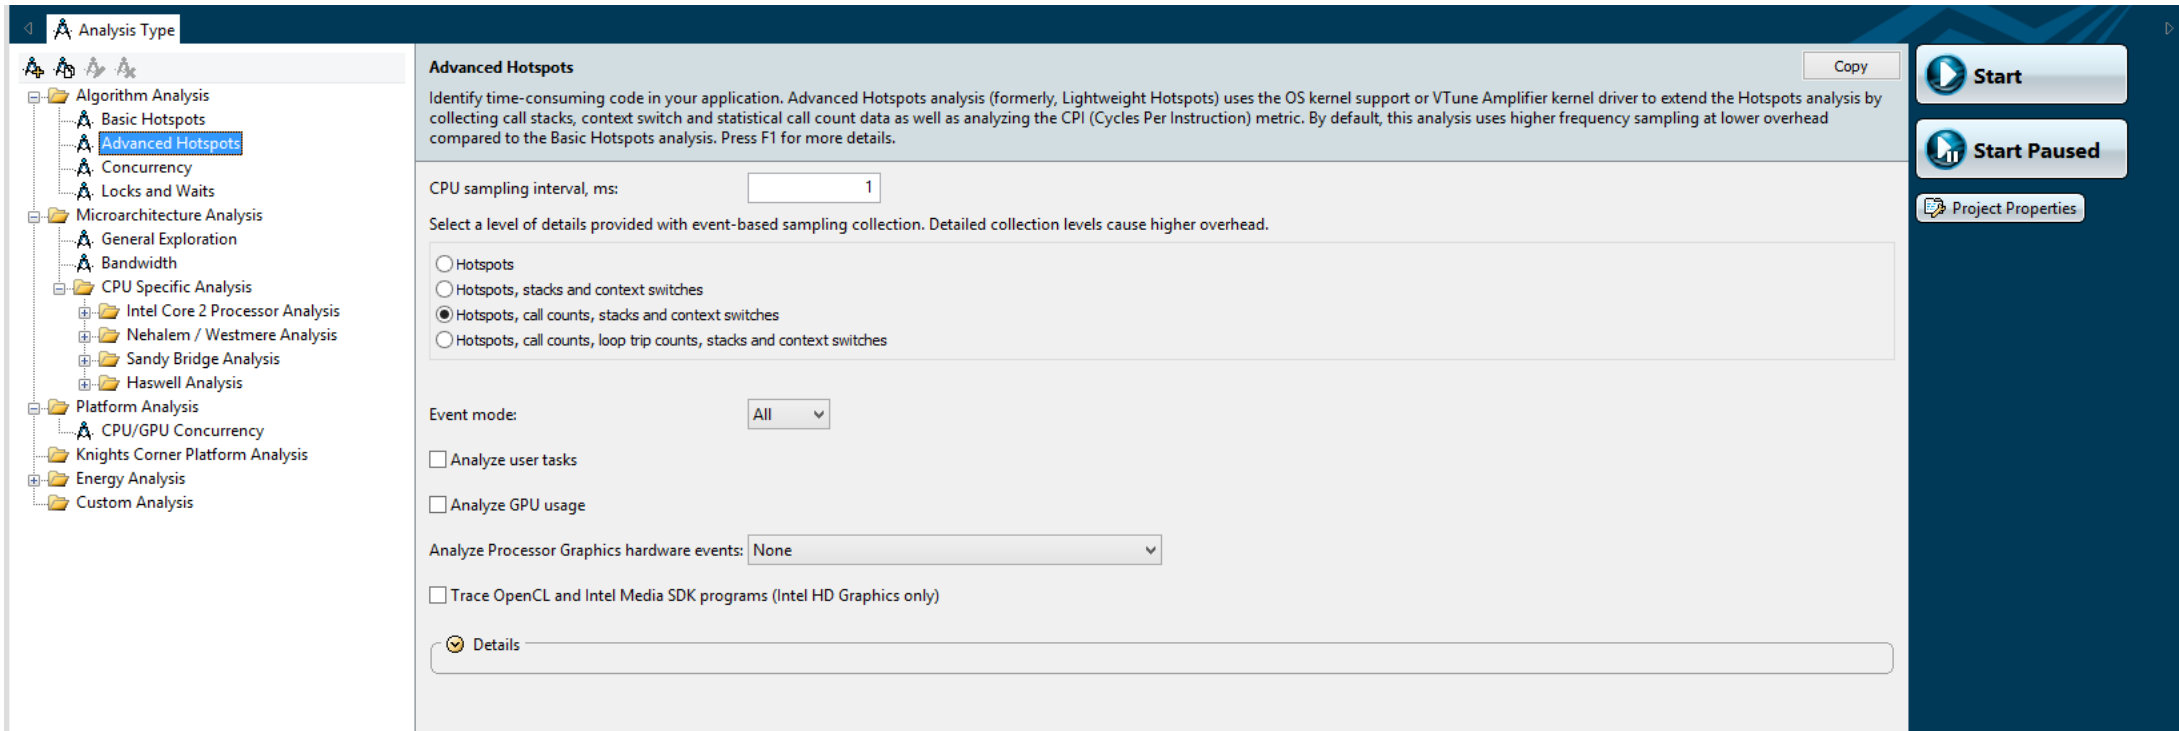
VTune Amplifier

- Helps in collections of performance data of an application running on the system
- Organizes data in different views
- Helps to identify potential issues and also suggests improvements

Vtune Example – Step 1: Create Project



Vtune Example – Step 2: Choose analysis type & start



VTune Example – Step 3: View Summary

⬆️ **Elapsed Time:** 53.202s 📄

⬆️ **CPU Time:** 48.692s

⬇️ **Effective Time:** 48.686s

Spin Time: 0.006s

Overhead Time: 0s

Instructions Retired: 41,371,428,877

Estimated Call Count: 46,382,814

CPI Rate: 2.065

The CPI may be too high. This could be caused by issues such as memory stalls, instruction starvation, branch misprediction or long latency instructions. Explore the other hardware-related metrics to identify what is causing high CPI.

Wait Rate: 5.262

CPU Frequency Ratio: 1.316

⬆️ **Context Switch Time:** 2.460s

Wait Time: 0.226s

Inactive Time: 2.234s

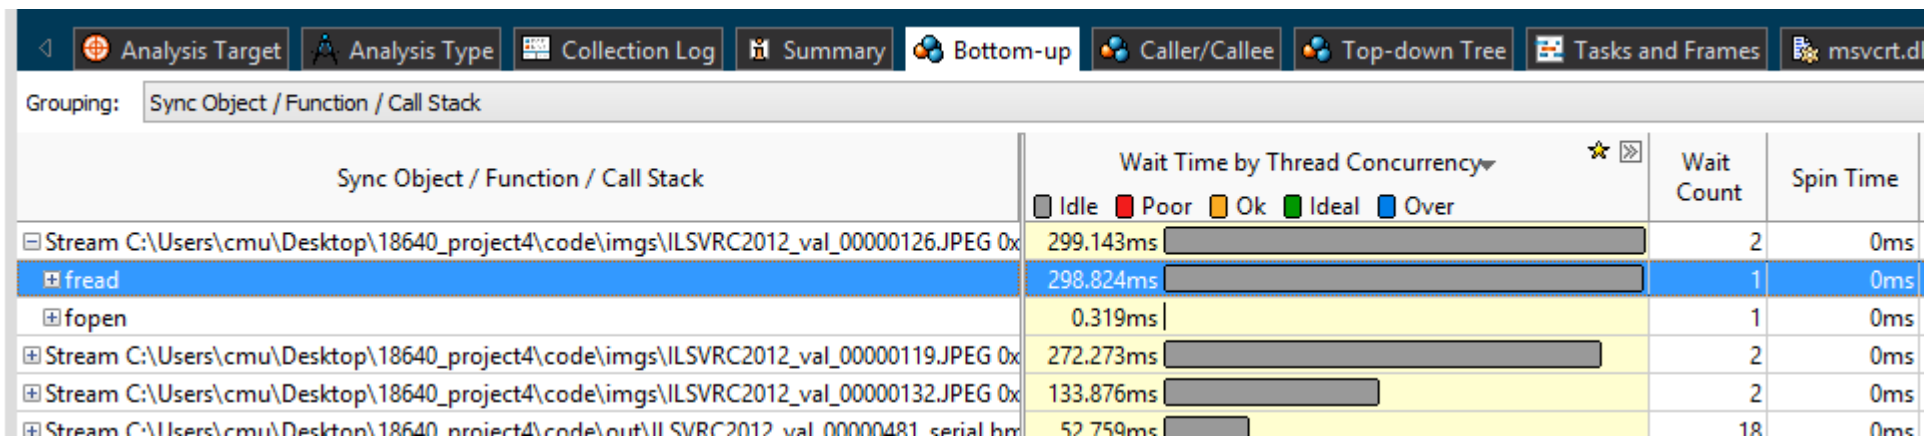
Paused Time: 0.116s

Estimated Call Count: 46,382,814

VTune Example – Some terms

- CPU time: time for which CPU is actively executing
- Effective time: time spent in the user code (excludes spin and overhead time)
- Wait time: time while threads were waiting for synchronization
- Cycles per instruction (CPI): average amount of time in machine clock cycles for each instruction

Vtune Example – Step 4: Analyze other tabs



Sync Object / Function / Call Stack					Wait Count	Spin Time
Wait Time by Thread Concurrency						
Legend: Idle (grey), Poor (red), Ok (orange), Ideal (green), Over (blue)						
Stream C:\Users\cmu\Desktop\18640_project4\code\imgs\ILSVRC2012_val_00000126.JPEG 0x	299.143ms				2	0ms
fread	298.824ms				1	0ms
fopen	0.319ms				1	0ms
Stream C:\Users\cmu\Desktop\18640_project4\code\imgs\ILSVRC2012_val_00000119.JPEG 0x	272.273ms				2	0ms
Stream C:\Users\cmu\Desktop\18640_project4\code\imgs\ILSVRC2012_val_00000132.JPEG 0x	133.876ms				2	0ms
Stream C:\Users\cmu\Desktop\18640_project4\code\out\ILSVRC2012_val_00000481_serial_b...	52.759ms				18	0ms

1. Click on function to get source code
2. From source code, one can also navigate to Instruction Manual (do install Adobe Reader)

Outline

- What is a compute stick?
- How to use compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

Digit Detection

- Collection of images where each image represents a digit between 0-9
- Every digit is represented by a 28x28 matrix
- Every cell value of the matrix represents a pixel with a discrete value between 0 - 255

Digit Detection Algorithm – Nearest Neighbor

1. Store a set of images (matrix and its label) in memory – *training set*
2. For every new image i in the *test set*:
 - I. find nearest neighbor by calculating the euclidean distance from every image in the training set
 - II. assign label of the training image which has the minimum distance to test image i

Outline

- What is a compute stick?
- How to use compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

Limitations of Compute Stick

- Not ideal for tasks which are computationally heavy
- Low memory so Vtune may crash
- Some restriction on collection of GPU metrics due to BIOS
- RDP not possible with the pre-installed version of Windows (alternative TightVNC)

Outline

- What is a compute stick?
- How to use compute stick?
- Multi-Core Execution Libraries
- Image Processing Library & Edge Detection Example
- Performance Evaluation using Intel's VTune
- Digit Detection
- Limitations of Compute Stick
- What you have to do?
- References

What you have to do?

- Set up Windows on Compute stick
 - Set Windows username as **cmu_18640**
 - Set password as **root**
- Install the following:
 - MinGW (Minimalist GNU for Windows)
 - OpenCL
 - Intel VTune Amplifier

What you have to do?

- Add proper paths to the system path variables. For example:
 - If you run multi-threaded with SSE (mt_sse) make it has in it's path the muti-threaded with SSE libraries in its path
- Compile the program using ming32-make
- Run VTune (**as administrator**) with each executable and observe the reading for different types of analysis

Report

- Answers all questions under Report sections in the handout.

Matrix Multiplication	Edge Detection	Digit Detect
mmm_single_thread_scalar	st_nsse	st_digit
mmm_single_thread_simd	st_sse	
mmm_multi_thread_scalar	mt_nsse	
mmm_multi_thread_simd	mt_sse	
	mt_sse_ocl	

Yes, we have a bonus section too!

- Parallelize the `st_digit` (digit detection) code to make it run faster.
- Go through the sequential code and see how you can parallelize it.
- Use any method to make it parallel.
- Report the same parameters for the parallel version as asked in previous part.
- Score bonus 10 points!

Please return compute sticks in its original box no later than final exam day

References

- Lecture notes by Simon Lucey (http://16423.courses.cs.cmu.edu/slides/Lecture_2.pdf)
- OpenMP: A short introduction, 18-740, CMU Fall 2009 (http://www.cs.cmu.edu/afs/cs/academic/class/15740-f09/public/asst/asst2/openmp_tutorial.pdf)
- OpenCL Overview (<http://sa10.idav.ucdavis.edu/docs/sa10-dg-opencl-overview.pdf>)
- Mohamed Zahran, NYU (<http://cs.nyu.edu/courses/fall12/CSCI-GA.3033-012/lecture7.pdf>)
- LLNL pages for OpenMP and OpenCL
- <https://www.khronos.org/opencl/>

Questions